

(L1) CRR Pricing Engine

(L2) ML Extensions

(L3) Cross-Model Comparison

I What Is the Problem?

Research Questions

- Do CRR model prices predict whether an equity option expires ITM or OTM?
- Does cross-validation between CRR and an ML classifier detect exploitable pricing *edge*?
- Does regime-conditioned Kelly sizing convert that edge into positive expected value?

Core Inefficiency Gap

Retail traders lack a principled framework to decide *whether* a price is fair, *how much* to bet, or *when not to trade*.

The Pricing Gap

The spread between CRR model price and market price — confirmed by the ML classifier, this gap signals actionable *edge*.

The Pricing Gap

Market: supply-demand noise, spread, and sentiment.

Model: pure risk-neutral expectation.

Gap = mispricing *edge* — only when both layers agree.

No quantified edge · no position-sizing rule
· no regime filter.

$$\text{Edge} = \frac{V_{\text{model}} - V_{\text{market}}}{V_{\text{market}}}$$

I Why It Matters & What Others Have Done

Why This Problem Matters

- Options trade **\$600B+ notional daily**; even 2% mispricing detection has large profit-and-loss impact
- **Fama EMH (1970)**: public info is priced in on average — but not in every regime
- **Herd**ing creates regime-local inefficiencies (Samuelson macro-inefficiency paradox)
- Retail traders lack tools to detect edge or size positions systematically


Prior Work

- **Cox, Ross & Rubinstein (1979)**: binomial lattice; no ML integration
- **Black & Scholes (1973)**: closed-form; fails for American early exercise
- **Bakshi et al. (1997)**: stochastic vol; fit requires unobservable parameters
- **Carr & Wu (2009)**: IV overstates realized vol (variance risk premium)
- **Kelly (1956)**: log-utility bet sizing; applied to markets by Thorp (1969)
- **Gap**: No framework cross-validates CRR + ML with regime-conditioned Kelly sizing

Novel Contribution

- **First** to cross-validate CRR against an IV-free ML classifier (fully independent signals)
- Regime-conditioned Kelly sizing from IV/RV₃₀ ratio
- Walk-forward 60/20/20 split; zero lookahead bias
- Open-source Streamlit calculator with live chain

Three-Layer Pipeline



Data AAPL 2016–2020, 101,535 contracts (walk-forward); AMD April 2026 live chain (live validation).

Worked Example: Contract A — HERDING Setup

Real AAPL put (2020-03-16) · $\sigma = \text{RV30}$, not IV

Contract A · HERDING REGIME — CROWD OVERPRICES

CRR binomial tree will now be built for this real AAPL put option

Quote Date:	2020-03-16
Underlying (S):	\$241.47
Strike (K):	\$240.00
DTE:	18 days
Time (T):	0.0493 years
RV30 (σ):	0.8224 ← used as sigma in CRR
IV:	1.0043
V_market:	\$21.225 (bid-ask mid)
Risk-free r:	2.00% (constant 2016-2020 avg)
CRR steps:	50 (pricing) / 5 (visualization)

Phase 0 of 4 | Next: build stock price lattice (forward pass)

Why this slide matters

Grounds the CRR pipeline in a real AAPL put pulled unmodified from the 2016–2020 Kaggle dataset.


Key reading

IV $\approx$ 1.00 vs RV30 $\approx$ 0.82 — the market pays for 22% more volatility than recent history justifies. Classic HERDING regime signature; entry point for the Layer 1 edge signal.

Setup snapshot: $S = \$241.47$, $K = \$240$, DTE = 18d, RV30 = 0.8224, IV = 1.0043, $V_{\text{market}} = \$21.225$, $r = 2.00\%$, $N = 50$ pricing steps.

Forward Pass — Building the Stock Price Lattice

u, d, p^* derived from $\sigma = \text{RV30}$ · 5-step visualization, 50-step pricing



What this pass does

Lays out every possible stock price path.
No option value computed yet — only S evolution from $t = 0$ to expiry.


What “no-arbitrage” means

No-arbitrage = no risk-free profit can exist in an efficient market.

Reading: Lattice recombines — node count grows linearly ($N+1$), not exponentially (2^N). That is why 50 pricing steps remain cheap.

Backward Induction — Where the Option Price Emerges

CRR fair value \$16.663 vs market \$21.225 · overpriced +21.5%



At every interior node

$$\text{intrinsic} = \max(K - S, 0)$$

$$\text{hold} = e^{-r\Delta t} (p \cdot V_{\text{up}} + q \cdot V_{\text{dn}})$$

$$V = \max(\text{intrinsic}, \text{hold})$$

Result at the root

$V(N=50) = \$16.663$ (CRR fair value)


$V_{\text{market}} = \$21.225$ (bid-ask mid)

Edge = -21.5% ★ OVERPRICED

Reading: When the gold ★ sits above root V, the crowd pays more than no-arbitrage theory permits. That gap is the Layer 1 edge — already negative here, before any ML model is consulted.

Contract A Result — Two Independent Models Agree

Layer 1 (CRR) + Layer 2 (XGBoost) → BOTH SELL · HERDING regime



Layer 1 (CRR math)

$$V_{\text{market}} = \$21.225 \quad V_{\text{model}} = \$16.663$$

Edge = -0.2149 → OVERPRICES

Layer 2 (XGBoost ML)

$$p_{\text{indep}} = 0.0562 \quad q_{\text{market}} = 0.0884$$


Edge = -0.0322 → OVERPRICES

BOTH SELL — HERDING regime

Why it matters: Neither layer was trained on the other. Agreement comes from shared underlying reality — the empirical fingerprint of Samuelson macro-inefficiency.


CRR Pricing Results

Model vs. Market




Each point = one option contract.
Y-axis: CRR model price; X-axis: market mid (101k AAPL contracts).
Tight 45° clustering confirms well-calibrated ITM/ATM pricing; deep-OTM put residuals expose the volatility risk premium.

Price Convergence (N -step)



CRR option price vs. tree depth N .
Characteristic odd-even oscillation damps rapidly; residual error drops below \$0.01 by $N \approx 100$. $N = 200$ used throughout to guarantee convergence within typical bid-ask spreads.

Early Exercise Boundary




Critical stock price $S^*(t)$ below which immediate exercise of an American put dominates holding. The boundary rises monotonically toward the strike as $T \rightarrow 0$, matching Brennan-Schwartz theory; confirms lattice captures American optionality.

AMD Live Validation: AAPL $S_0 = \$275.10$, $K = \$280$, $\sigma = 64\%$, $T = 18d$ — model within 1.2% of market mid on 80% of contracts.


Cross-Validation, Regime Detection & Backtest

L1 vs. L2 Cross-Validation


Layer 1 (CRR) vs Layer 2 (XGBoost) — Cross-Model Comparison
AAPL Options 2016-2020 - Same 3,000 Contracts



Regime Analysis (IV / RV₃₀)



Monte Carlo Profit and Loss (1,000 Trades)



Key Findings

- **L1 vs L2:** when both signals say SELL (Q3), actual ITM rate jumps to **76.5%** vs. ~63% baseline — evidence the cross-validation has predictive content.
- **Regime Analysis:** IV/RV₃₀ splits market; aggregate $r = 0.032$ weak, herding $r = 0.413$ ($n=217$, $p=2.4 \times 10^{-10}$) — Simpson's Paradox masks the signal in aggregate (Samuelson).
- **Profit and Loss:** \$2.9M cumulative on \$100k base; max draw-down < 15%, Brier=0.211; regime-conditioned Kelly (half-Normal quarter-Herding) converts edge to positive EV.

Research Extensions

- 1 Expand universe to SPY / QQQ / NDX options for a broader cross-sectional market test
- 2 Integrate LSTM volatility forecaster to replace static IV inputs with dynamic IV surface predictions
- 3 Deploy live paper-trading loop via broker API (IBKR / Alpaca) with real-time signal generation
- 4 Extend pricer to multi-leg spreads — vertical, calendar, iron condor — with joint Kelly sizing
- 5 Publish open-source framework and Streamlit calculator for community validation and peer review

Known Limitations

Data scope: AAPL 2016–2020 is a single equity in a single market cycle; out-of-sample performance on other tickers is untested.

Volatility surface: CRR uses a flat IV input per contract; no smile or skew interpolation across strikes.

Transaction costs: Backtest assumes mid-price fills; bid-ask friction and slippage are not modeled.

Regime detection lag: IV/RV₃₀ ratio is backward-looking; regime transitions are identified only after the fact.

Kelly assumptions: Log-utility maximization assumes no leverage constraints or portfolio-level risk limits.

Live Demo — Interactive CRR Binomial Pricing Calculator

Streamlit application · runs locally on any laptop with Python 3.13 · ticker-agnostic

Deploy

Contract Parameters

Ticker (for RV30 lookup)

AAPL

Option Type Action

put buy

Underlying Price (\$5)

200.00

Strike Price (K \$)

200.00

Market Price (V_{market})

10.00

Implied Volatility (IV)

0.3000

CRR Binomial Pricing Calculator

Cox-Ross-Rubinstein American option pricing · Kelly Criterion sizing · Monte Carlo simulation | Works for any ticker · No dividends assumed

[Pricing](#) [Greeks](#) [Edge & Signal](#) [Kelly Sizing](#) [Monte Carlo](#) [Tree Animation](#) [User's Guide](#)

Enter your contract parameters in the sidebar and click Calculate to populate this tab.

Demo plan (next 3 minutes)

- Reproduce Contract A end-to-end
- Confirm L1 edge $\approx -21.5\%$ (SELL)
- Show Kelly sizing on \$100k capital
- Animate the binomial tree live
- Live RV30 lookup for any ticker

I Conclusion

Central Contribution

CRR (math) and XGBoost (ML) — using **independent inputs** — both agree on mispricing direction in the herding regime, mutually cross-validating Samuelson macro-inefficiency (Fama, 1970).

Hypotheses Confirmed

H1 — CRR Benchmark Confirmed. Model within 1–2% of market mid on 80% of AMD live-chain contracts; sub-rounding precision validates CRR as fair-value reference.

H2 — XGBoost Independent Estimator Confirmed. Brier score $0.211 < 0.25$ naive baseline; top features are moneyness and Days to Expiration (DTE) — no implicit IV leakage.

H3 — Fractional Kelly Dominates Confirmed. \$2.9M P&L over 101,535 AAPL contracts (2016–2020); max drawdown $< 15\%$, circuit breaker never triggered.

Practical Implication: The pipeline's most valuable output is the *no-trade signal* — when CRR and XGBoost disagree, no position is taken. The model works precisely *because* it says no when there is no edge.